

Why Vulnerabilities are Hard to Eliminate

UT CS361S

FALL 2021

LECTURE NOTES

A solid orange horizontal bar at the bottom of the slide.



Make it harder to control
a subverted flow




Make taking control of
the flow innocuous

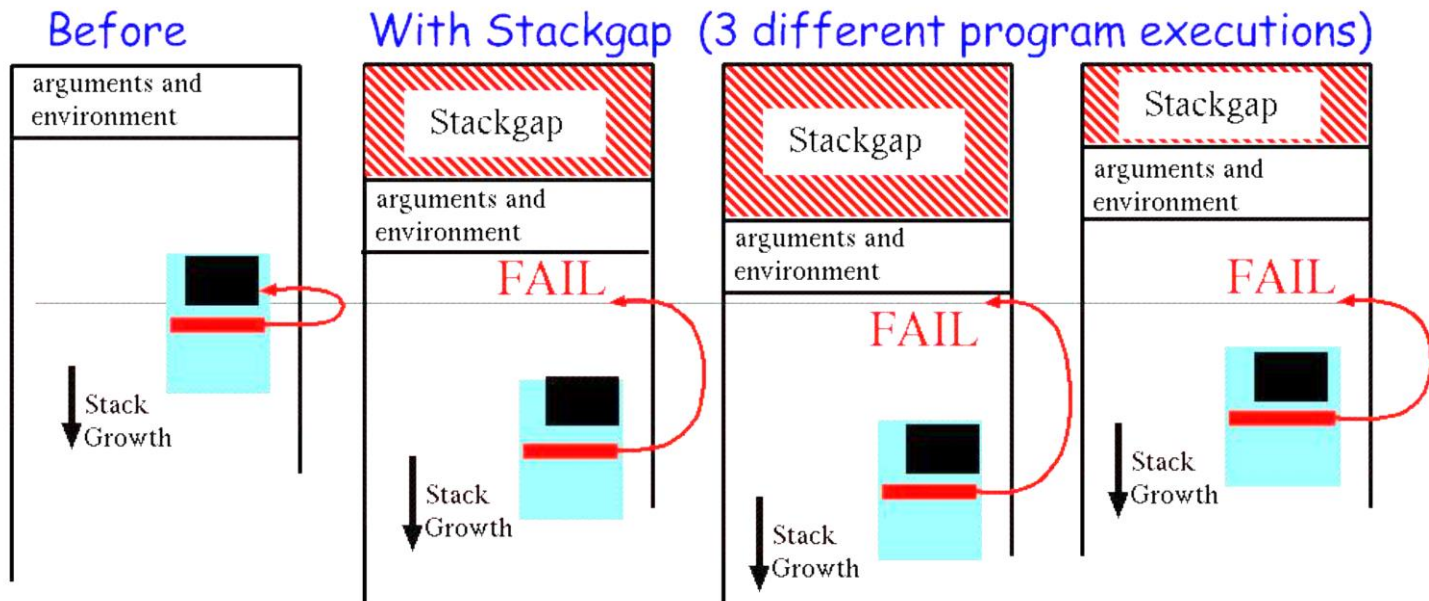


Make it harder to get
control of the flow

Disrupting Exploitative Operations



Random Stack Gap



ASLR

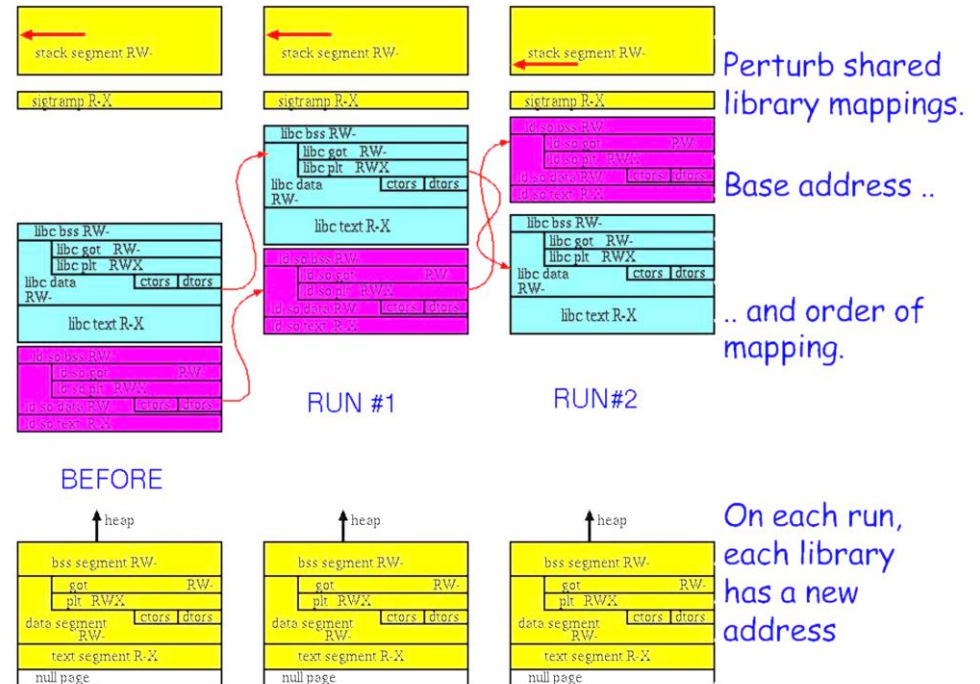
Address Space Layout Randomization

Subversion usually needs to know memory layout

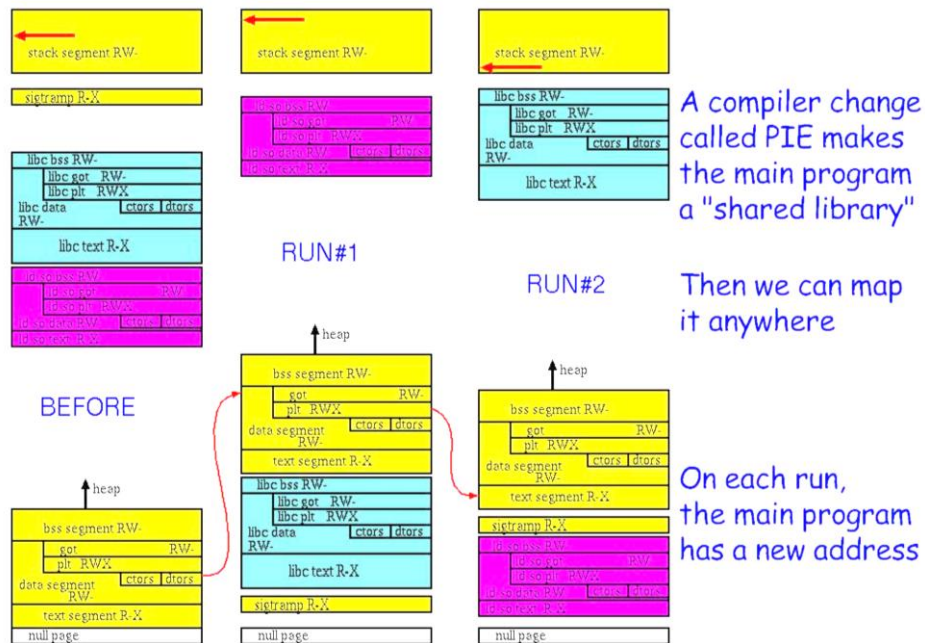
General goal: make layout unpredictable

Start With Libraries

ASLR: randomly map & order libraries



PIE - Position Independent Executable



Add Executables

Finally, Dynamic Allocations



The diagram consists of two identical graphic elements side-by-side. Each element features a solid orange rounded rectangle in the background. Overlapping the bottom-right portion of this orange rectangle is a white rounded rectangle with a thin orange border. The text 'mmap' is centered within the white rectangle.

mmap



The diagram consists of two identical graphic elements side-by-side. Each element features a solid orange rounded rectangle in the background. Overlapping the bottom-right portion of this orange rectangle is a white rounded rectangle with a thin orange border. The text 'malloc' is centered within the white rectangle.

malloc

Limitations of ASLR

1. **Boot-time based randomization**
2. **Unsupported executables/libraries, low-entropy.**
3. **ASLR does not *trap* the attack**
4. **ASLR does not alert in a case of an attack**
5. **ASLR does not *provide information* about an attack**
6. **ASLR is being bypassed by exploits daily**

Posted by **MORDECHAI GURI, PH.D.** on December 17, 2015

Making Violations Less Dangerous



W^X
Permissions

The diagram consists of two overlapping rounded rectangles. The background rectangle is orange, and the foreground rectangle is light orange. The text W^X and Permissions is centered within the light orange rectangle.




rodata

The diagram consists of two overlapping rounded rectangles. The background rectangle is orange, and the foreground rectangle is light orange. The text rodata is centered within the light orange rectangle.

W | x Permissions

Many bugs are exploitable because the address space has memory that is both writeable and executable (permissions = W | X)



this location has to be executable for the exploit to work

We could make the stack non-executable...

Hmmmm... how about a generic policy for the whole address space:

A page may be either writeable or executable, but not both (unless the program specifically requests)

We call this policy $W \wedge X$ ($W \text{ xor } X$)

Executable Stacks

This is what static executables
used to look like in memory.

The stack has a piece of executable
called the "signal trampoline"

What is this?



First problem: The stack is
executable

5.6 Returning from a signal handler

When the program was interrupted by a signal, its status (including all integer and floating point registers) was saved, to be restored just before execution continues at the point of interruption.

This means that the return from the signal handler is more complicated than an arbitrary procedure return - the saved state must be restored.

To this end, the kernel arranges that the return from the signal handler causes a jump to a short code sequence (sometimes called *trampoline*) that executes a `sigreturn()` system call. This system call takes care of everything.

In the old days the trampoline lived on the stack, but nowadays (since 2.5.69) we have a trampoline in the [vsyscall](#) page, so that this trampoline no longer is an obstacle in case one wants a non-executable stack.

Linux Trampoline?

Linux Trampoline!!!

No-execute stacks [\[edit \]](#)

Some implementations of trampolines cause a loss of [no-execute stacks](#) (NX stack). In the [GNU Compiler Collection](#) (GCC) in particular, a nested function builds a trampoline on the stack at runtime, and then calls the nested function through the data on stack. The trampoline requires the stack to be executable.

No execute stacks and nested functions are mutually exclusive under GCC. If a nested function is used in the development of a program, then the NX stack is silently lost. GCC offers the `-Wtrampoline` warning to alert of the condition.

Software engineered using [secure development lifecycle](#) often do not allow the use of nested functions due to the loss of NX stacks.^[11]

[.wikipedia.org/wiki/Trampoline_\(computing\)#No-execute_stacks](https://en.wikipedia.org/wiki/Trampoline_(computing)#No-execute_stacks)

The .rodata Segment

W^X Transition: The .rodata segment

Readonly strings and pointers were stored in the `.text` segment: X | R

Meaning const data could be executed (could be code an attacker could use as ROP payload)


Solution: start using the ELF `.rodata` segment

These objects are now only R, lost their X permission

Greater policy: "minimal set of permissions"

Finally, Blocking Exploits

Stack Protector



A typical stack frame...

Random value is inserted here by function prologue ...
... and checked by function epilogue

Reordering: Arrays (strings) placed closer to random value -- integers and pointers placed further away

`-fstack-protector-all` compiled system is 1.3% slower at make build

return-Oriented PROGRAMMING


David Brumley

Carnegie Mellon University

Credit: Some slides from Ed Schwartz

Control Flow Hijack:

Always control + computation




```
"\x31\x09\xf7\xe1\x51\x68\x2f\x2f"  
"\x73\x68\x68\x2f\x62\x69\x6e\x89"  
"\xe3\xb0\x0b\xcd\x80";
```

Previously: Executable code as input

Control Flow Hijack:

Always control + computation



Today: Return Oriented Programming
Execution without injecting code

ROP Overview

Idea:

We forge shell code out of existing application logic gadgets

Requirements:

vulnerability + gadgets + some unrandomized code
(we need to know the addresses of gadgets)



Technically, *PREDICTABLE*

Motivation: Return-to-libc Attack


ret transfers control to system,
which finds arguments on stack

Overwrite return address with address of libc
function

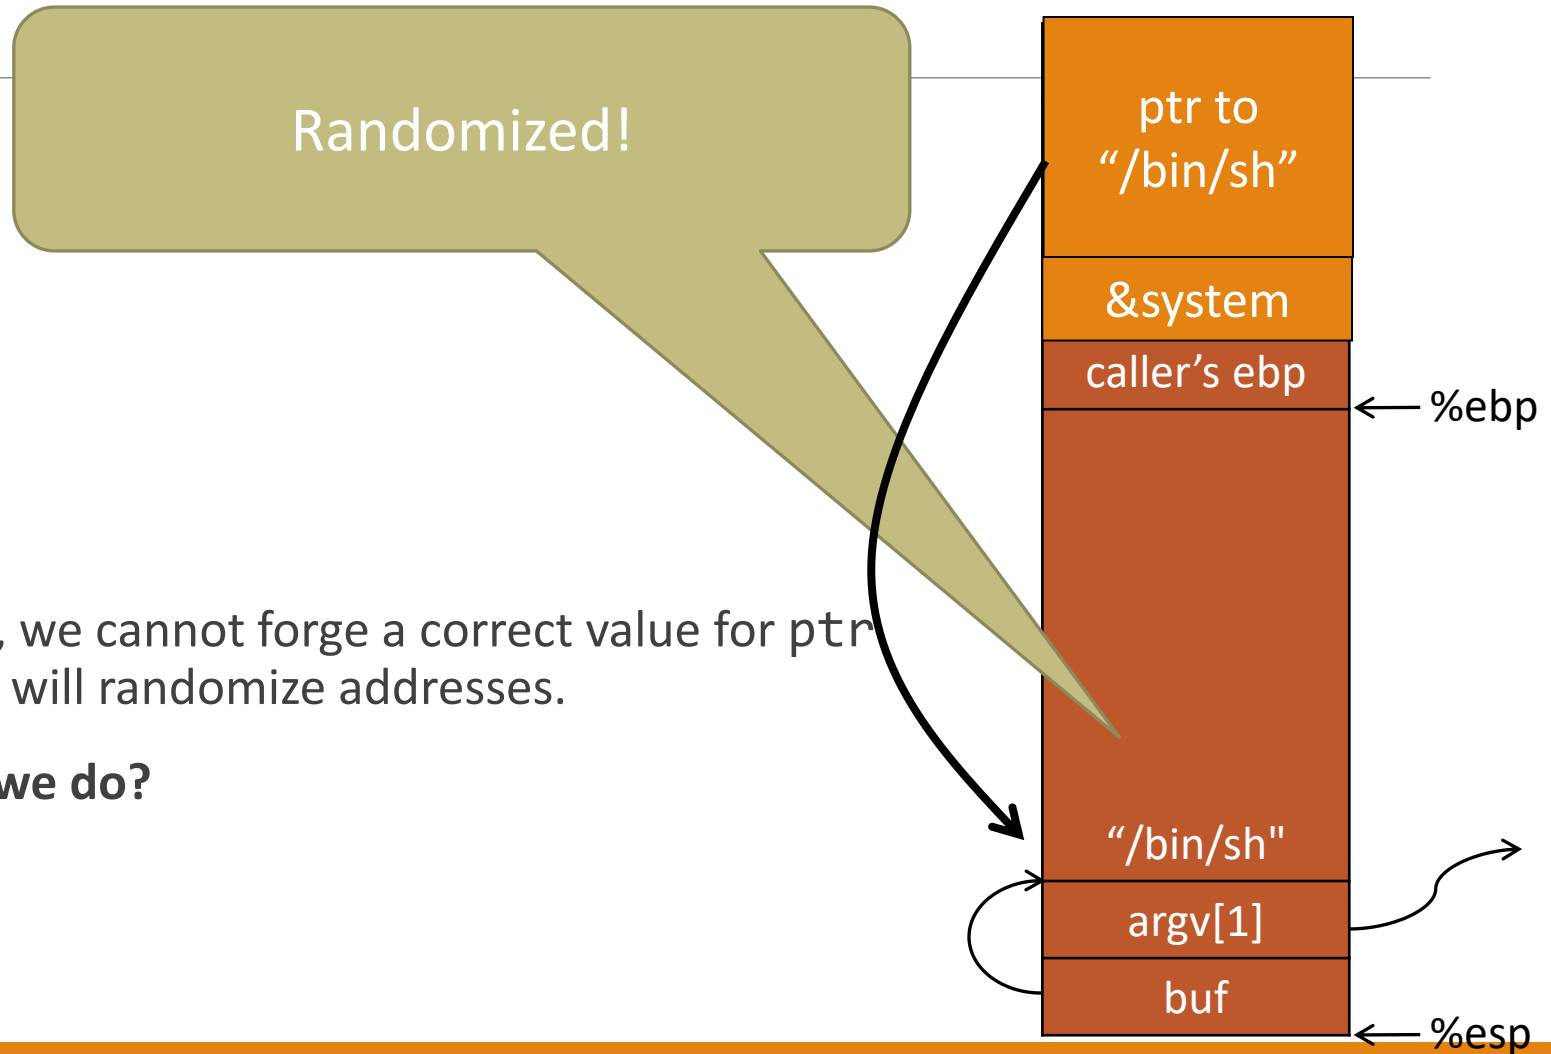
setup fake return address and argument(s)

ret will "call" libc function

No injected code!



Question



With ASLR, we cannot forge a correct value for ptr since ASLR will randomize addresses.

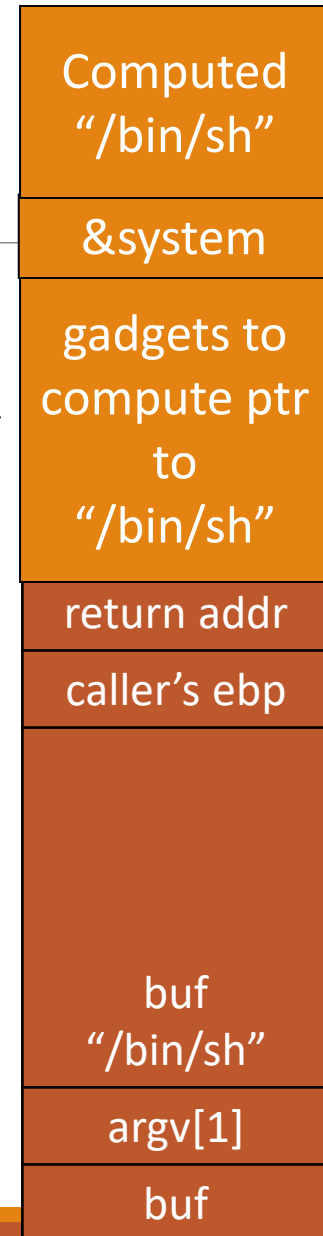

What can we do?

Idea!

Get a copy of ESP to calculate address of
“/bin/sh” on randomized stack.

This works because ASLR only protects against knowing *absolute* addresses, while we will find it's *relative address*.

Writes



Return Chaining

Suppose we want to call 2 functions in our exploit:

foo(arg1, arg2)

bar(arg3, arg4)

Stack unwinds up

First function returns into code to advance stack pointer

- e.g., pop; pop; ret

What does this do?

Overwritten
ret addr

arg4
arg3
&(pop-pop-ret)
bar
arg2
arg1
&(pop-pop-ret)
foo

Return Chaining

When **foo** is executing, &pop-pop-ret is at the saved EIP slot.

When **foo** returns, it executes pop-pop-ret to clear up arg1 (pop), arg2 (pop), and transfer control to **bar** (ret)

arg4
arg3
&(pop-pop-ret)
bar
arg2
arg1
&(pop-pop-ret)
foo


There are many
semantically equivalent
ways to achieve the
same net shellcode
effect

Let's practice thinking in gadgets

An example operation

Mem[v2] = v1

Desired Logic



a_1 : mov eax, [esp] ; eax has v1

a_2 : mov ebx, [esp+8] ; ebx has v2

a_3 : mov [ebx], eax ; Mem[v2] = eax

Implementation 1

implementing with gadgets


Mem[v2] = v1

Desired Logic

eax	v_1
ebx	
eip	a_1

a_1 : pop eax
 a_2 : ret
 a_3 : pop ebx
 a_4 : ret
 a_5 : mov [ebx], eax

Implementation 2




Suppose a_5
and a_3 on
stack

implementing with gadgets

Mem[v2] = v1

Desired Logic

eax	v ₁
ebx	
eip	a ₃



Stack

a₁: pop eax
a₂: **ret**
a₃: pop ebx
a₄: ret
a₅: mov [ebx], eax


Implementation 2

implementing with gadgets

Mem[v2] = v1

Desired Logic

eax	v ₁
ebx	v ₂
eip	a ₃



Stack

a₁: pop eax
a₂: ret
a₃: pop ebx
a₄: ret
a₅: mov [ebx], eax


Implementation 2

implementing with gadgets

Mem[v2] = v1

Desired Logic

eax	v ₁
ebx	v ₂
eip	a ₅



Stack

```
a1: pop eax;  
a2: ret  
a3: pop ebx;  
a4: ret  
a5: mov [ebx], eax
```


Implementation 2

implementing with gadgets

Mem[v2] = v1

Desired Logic

eax	v ₁
ebx	v ₂
eip	a ₅



Stack

```
a1: pop  eax; } Gadget 1
a2: ret      }
a3: pop  ebx; } Gadget 2
a4: ret      }
a5: mov  [ebx], eax
```


Implementation 2

Equivalence

Mem[v2] = v1

Desired Logic

semantically
equivalent



Stack

“Gadgets”

↔ a₁: pop eax; ret
↔ a₂: pop ebx; ret
↔ a₃: mov [ebx], eax

Implementation 2


Return-Oriented Programming (ROP)

Mem[v2] = v1

Desired *Shellcode*

a₁: pop eax; ret
a₂: pop ebx; ret
a₃: mov [ebx], eax

Desired store executed!



Gadgets

A gadget is a set of instructions for carrying out a semantic action

- mov, add, etc.

Gadgets typically have a number of instructions

- One instruction = native instruction set
- More instructions = synthesize <- ROP

Gadgets in ROP generally (but not always) end in return

Return-Oriented Programming

is A lot like a ransom
note, BUT instead of cutting
cut letters from magazines,
YOU ARE cutting out
instructions from text
segments

RO(P?) Programming

1. Disassemble code
2. Identify useful code sequences as gadgets
3. Assemble gadgets into desired shellcode

Attacker Oriented Programming?

Behavior isn't a program

We should be able to perfectly detect bad behavior, right?


“Weird Machines”

“Weird machines, exploitability, and provable unexploitability”

Written by Thomas Dullien

Explains that users interacting with a program *is a program*

What is a Program?




*From Dullien's Paper

State Machine View

View a “Program” as a state machine

Program starts in state S_0

Based on instruction, advances to state S_i




States and User Interactions

Program is in some State. Call it S_0

User interacts with the program

Program advances to state S_1



What is a “User”?

Do we literally mean a flesh-and-blood human?

Really, “user” is just whatever provides the input

This can, of course, just be another process

Thus, two processes interacting ***IS A PROGRAM***

Therefore, determining if “behavior” is good is undecidable

One More Big Problem

Decidability is a fundamental, ***unsolvable*** problem

Another big problem is ***Supply Chain***

1984: Thompson's Reflections

“Reflections on Trusting Trust” by Ken Thompson, 1984

Demonstrated creating an evil compiler

Would compile a login program with a backdoor

BUT! ***ALSO COMPILED COMPILERS WITH THIS LOGIC!***

“Clean” compiler source code compiled by an evil compiler ***is evil!***

Proved that a “source code review” can’t catch all evil